Fundamentals of Artificial Intelligence and Knowledge Representation

Contents

1 Introduction

1.1 AI systems classification

1.1.1 Intelligence classification

Intelligence is defined as the ability to perceive or infer information and to retain the knowledge for future use.

Weak AI aims to build a system that acts as an intelligent system.

<div style="text-align: right">Weak AI</div>

Strong AI aims to build a system that is actually intelligent.

<div style="text-align: right">Strong AI</div>

1.1.2 Capability classification

General AI systems able to solve any generalized task.

<div style="text-align: right">General AI</div>

Narrow AI systems able to solve a particular task.

<div style="text-align: right">Narrow AI</div>

1.1.3 AI approaches

Symbolic AI (top-down) Symbolic representation of knowledge, understandable by humans.

<div style="text-align: right">Symbolic AI</div>

Connectionist approach (bottom up) Neural networks. Knowledge is encoded and not understandable by humans.

<div style="text-align: right">Connectionist approach</div>

1.2 Symbolic AI

Deductive reasoning Conclude something given some premises (general to specific). It is unable to produce new knowledge.

<div style="text-align: right">Deductive reasoning</div>

> **Example.** "All men are mortal" and "Socrates is a man" → "Socrates is mortal"

Inductive reasoning A conclusion is derived from an observation (specific to general). Produces new knowledge, but correctness is not guaranteed.

<div style="text-align: right">Inductive reasoning</div>

> **Example.** "Several birds fly" → "All birds fly"

Abduction reasoning An explanation of the conclusion is found from known premises. Differently from inductive reasoning, it does not search for a general rule. Produces new knowledge, but correctness is not guaranteed.

<div style="text-align: right">Abduction reasoning</div>

> **Example.** "Socrates is dead" (conclusion) and "All men are mortal" (knowledge) → "Socrates is a man"

Reasoning by analogy Principle of similarity (e.g. k-nearest-neighbor algorithm).

<div style="text-align: right">Reasoning by analogy</div>

> **Example.** "Socrates loves philosophy" and Socrates resembles John → "John loves philosophy"

Constraint reasoning and optimization Constraints, probability, statistics.

<div style="text-align: right">Constraint reasoning</div>

1.3 Machine learning

1.3.1 Training approach

Supervised learning Trained on labeled data (ground truth is known). Suitable for classification and regression tasks.

Unsupervised learning Trained on unlabeled data (the system makes its own discoveries). Suitable for clustering and data mining.

Semi-supervised learning The system is first trained to synthesize data in an unsupervised manner, followed by a supervised phase.

Reinforcement learning An agent learns by simulating actions in an environment with rewards and punishments depending on its choices.

1.3.2 Tasks

Classification Supervised task that, given the input variables X and the output (discrete) categories Y, aims to approximate a mapping function $f : X \to Y$.

Regression Supervised task that, given the input variables X and the output (continuous) variables Y, aims to approximate a mapping function $f : X \to Y$.

Clustering Unsupervised task that aims to organize objects into groups.

1.3.3 Neural networks

A neuron (**perceptron**) computes a weighted sum of its inputs and passes the result to an activation function to produce the output.

Figure 1.1: Representation of an artificial neuron

A **feed-forward neural network** is composed of multiple layers of neurons, each connected to the next one. The first layer is the input layer, while the last is the output layer. Intermediate layers are hidden layers.

The expressivity of a neural networks increases when more neurons are used:

**Single perceptron** Able to compute a linear separation.

Figure 1.2: Separation performed by one perceptron

Three-layer network Able to separate a convex region ($n_\text{edges} \leq n_\text{hidden neurons}$)
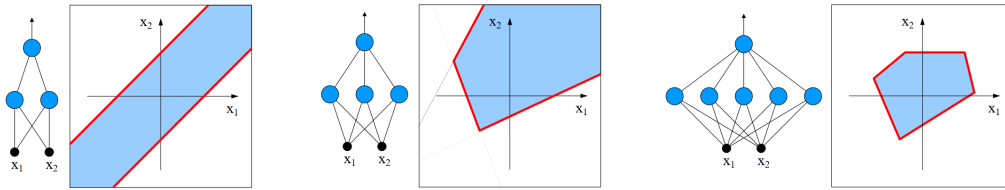

Figure 1.3: Separation performed by a three-layer network

Four-layer network Able to separate regions of arbitrary shape.
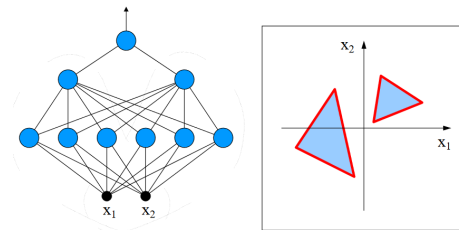

Figure 1.4: Separation performed by a four-layer network

Theorem 1.3.1 (Universal approximation theorem)**.** A feed-forward network with one hidden layer and a finite number of neurons is able to approximate any continuous function with desired accuracy.

Deep learning Neural network with a large number of layers and neurons. The learning process is hierarchical: the network exploits simple features in the first layers and synthesis more complex concepts while advancing through the layers.

1.4 Automated planning

Given an initial state, a set of actions and a goal, **automated planning** aims to find a partially or totally ordered sequence of actions to achieve a goal.
An **automated planner** is an agent that operates in a given domain described by:

- Representation of the initial state

- Representation of a goal

- Formal description of the possible actions (preconditions and effects)

1.5 Swarm intelligence

Decentralized and self-organized systems that result in emergent behaviors.

1.6 Decision support systems

Knowledge based system Use knowledge (and data) to support human decisions. Bottlenecked by knowledge acquisition.

Different levels of decision support exist:

Descriptive analytics Data are used to describe the system (e.g. dashboards, reports, . . .). Human intervention is required.

Diagnostic analytics Data are used to understand causes (e.g. fault diagnosis) Decisions are made by humans.

Predictive analytics Data are used to predict future evolutions of the system. Uses machine learning models or simulators (digital twins)

Prescriptive analytics Make decisions by finding the preferred scenario. Uses optimization systems, combinatorial solvers or logical solvers.

2 Search problems

2.1 Search strategies

Solution space Set of all the possible sequences of actions an agent may apply. Some of these lead to a solution.

Search algorithm Takes a problem as input and returns a sequence of actions that solves the problem (if exists).

2.1.1 Search tree

Expansion Starting from a state, apply a successor function and generate a new state.

Search strategy Choose which state to expand. Usually is implemented using a fringe that decides which is the next node to expand.

Search tree Tree structure to represent the expansion of all states starting from a root (i.e. the representation of the solution space).

Nodes are states and branches are actions. A leaf can be a state to expand, a solution or a dead-end. Algorithm 1 describes a generic tree search algorithm.
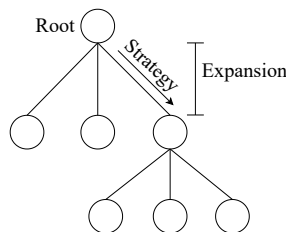

Figure 2.1: Search tree

Each node contains:

- The state
- The parent node
- The action that led to this node
- The depth of the node
- The cost of the path from the root to this node

2.1.2 Strategies

Non-informed strategy Domain knowledge not available. Usually does an exhaustive search.

Informed strategy Use domain knowledge by using heuristics.

Algorithm 1 Tree search algorithm

```python
def treeSearch(problem, fringe):
    fringe.push(problem.initial_state)
    # Get a node in the fringe and expand it if it is not a solution
    while fringe.notEmpty():
        node = fringe.pop()
        if problem.isGoal(node.state):
            return node.solution
        fringe.pushAll(expand(node, problem))
    return FAILURE

def expand(node, problem):
    successors = set()
    # List all neighboring nodes
    for action, result in problem.successor(node.state):
        s = new Node(
            parent=node, action=action, state=result, depth=node.dept+1,
            cost=node.cost + problem.pathCost(node, s, action)
        )
        successors.add(s)
    return successors
```

2.1.3 Evaluation

Completeness if the strategy is guaranteed to find a solution (when exists).

Time complexity time needed to complete the search.

Space complexity memory needed to complete the search.

Optimality if the strategy finds the best solution (when more solutions are possible).

2.2 Non-informed search

2.2.1 Breadth-first search (BFS)

Always expands the less deep node. The fringe is implemented as a queue (FIFO).

Completeness	Yes
Optimality	Only if cost is uniform (i.e. all edges have same cost)
Time and space complexity	$O(b^d)$, where the depth is d and the branching factor is b (i.e. each non-leaf node has b children)

The exponential space complexity makes BFS impractical for large problems.

2.2.2 Uniform-cost search

Same as BFS, but always expands the node with the lowest cumulative cost.

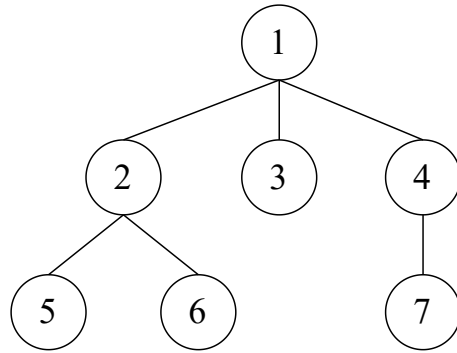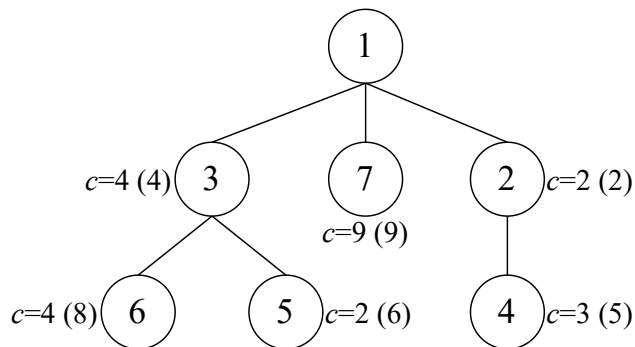| Completeness	Yes
Optimality	Yes
Time and space complexity	$O(b^d)$, with depth d and branching factor b

Figure 2.2: BFS visit order

Figure 2.3: Uniform-cost search visit order. (n) is the cumulative cost

2.2.3 Depth-first search

Always expands the deepest node. The fringe is implemented as a stack (LIFO).

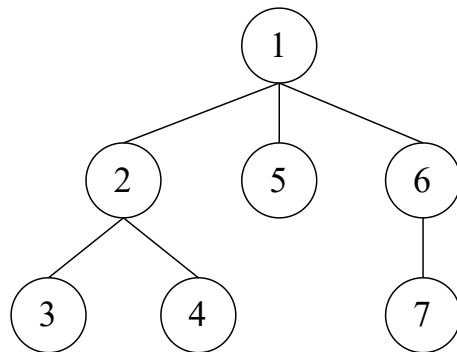| Completeness	No
Optimality	No
Time complexity	$O(b^d)$, with depth d and branching factor b
Space complexity	$O(b \cdot d)$, with depth d and branching factor b

Figure 2.4: DFS visit order