

Ethics in Artificial Intelligence (Module 3)

Last update: 14 March 2025

Academic Year 2024 – 2025
Alma Mater Studiorum · University of Bologna

Contents

- 1 Human agency and oversight 1**
 - 1.1 Governance and methodology 1
 - 1.2 HITL state-of-the-art approaches 2
 - 1.2.1 Active learning 2
 - 1.2.2 Interactive machine learning 2
 - 1.2.3 Machine teaching 2
- 2 Technical robustness and safety 4**
 - 2.1 Robust learning 5
 - 2.1.1 Robustness to model errors 6
 - 2.1.2 Robustness to unmodeled phenomena 6

1 Human agency and oversight

AI act, article 14 Article related to human oversight. It states that:

AI act, article 14

- Human centric AI is one of the key safeguarding principles to prevent risks.
- AI systems must be designed and developed with appropriate interfaces to allow humans to oversee them.

Human agency AI systems should empower human beings such that they can:

Human agency

- Make informed decisions.
- Foster their fundamental rights.

This can be achieved with methods like:

- Human-centric approaches,
- AI for social good,
- Human computation,
- Interactive machine learning.

Human oversight Oversight mechanisms to prevent manipulation, deception, conditioning from AI systems.

Human oversight

Possible methods are:

- Human-in-the-loop,
- Human-on-the-loop,
- Human-in-command.

Human-centered AI framework Approach centered on high autonomy while keeping human control.

Human-centered AI framework

Remark. Human agency and oversight happens at different levels:

Development team Responsible for the technical part.

Organization Decides who is in charge of accountability, validation, ...

External reviewers (e.g., certification entities).

1.1 Governance and methodology

Human-out-of-the-loop The environment is static and cannot integrate human knowledge. The AI system is a black-box that cannot be used in safety-critical settings.

Human-out-of-the-loop

Human-in-the-loop (HITL) The environment is dynamic and can use expert knowledge. The AI system is explainable or transparent and suitable for safety-critical settings.

Human-in-the-loop (HITL)

In practice, the AI system stops and waits for human commands before making a decision.

Society-in-the-loop	The society, with its conflicting interests and values, is taken into account.	Society-in-the-loop
Human-on-the-loop (HOTL)	The AI system operates autonomously and the human can intervene if needed.	Human-on-the-loop (HOTL)

Remark. Limitations of human-centric AI are:

- It does not scale well as human intervention is involved.
- It is hard to evaluate its effectiveness.
- Performance of the AI system might degrade.

1.2 HITL state-of-the-art approaches

1.2.1 Active learning

Active learning The system is in control of the learning process and the human acts as an oracle for labeling data. Active learning

The learner can query, following some strategy, the human for the ground-truth of unlabeled data. A general algorithm works as follows:

1. Split the data into an initial (small) pool of labeled data and a pool with the remaining unlabeled ones.
2. The model selects an example(s) to be labeled by the oracle.
3. The model is trained on the available labeled data.
4. Repeat until a stop condition is met.

The selection strategy can be:

Random

Uncertainty-based Select examples classified with the least confidence according to some metric.

Diversity-based Select examples that are rare or representative according to some metric.

Remark. This approach is effective in settings with lots of unlabeled data and annotating all of it is expensive.

Remark. This approach is sensitive to the choice of the oracle.

1.2.2 Interactive machine learning

Interactive machine learning Users interactively supply information that influences the learning process. Interactive machine learning

Remark. Compared to active learning, with interactive machine learning it is the human that selects the learning data.

1.2.3 Machine teaching

Machine teaching Human experts are completely in control of the learning process. There can be different types of teachers: Machine teaching

Omniscient teacher Complete access to the components of the learner (i.e., feature space, parameters, loss, optimization algorithm, ...).

Surrogate teacher Access to the loss.

Imitation teacher The teacher uses a copy of the learner that it can query to create a surrogate model.

Active teacher The teacher queries the learner and evaluates it based on the output.

Adaptive teacher The teacher selects examples based on the current hypothesis of the learner.

2 Technical robustness and safety

AI act, article 15 Article related to accuracy, robustness, and cybersecurity. It states that high-risk AI systems should:

- Be benchmarked and evaluated adequately.
- Be resilient to errors.
- Have measures to prevent and respond to attacks.

Technical robustness and safety AI systems should be secured to prevent unintentional harm and minimize the consequences of intentional harm. These requirements can be achieved by:

- Improving resilience to attacks.
- Introducing fallback plans.
- Improving general safety.
- Improving accuracy, reliability, and reproducibility.

Remark. Robustness is required as the real-world distribution is usually different from the training one.

Remark (Reliability vs robustness vs resilience).

Reliability Perform similarly on any test set from the same distribution.

In practice, reliable design aims at obtaining a probability of failure below some threshold.

Robustness Perform reasonably well on test sets from a slightly different distribution.

In practice, robust design aims at obtaining a model insensitive to small changes.

Resilience Adapt to unexpected inputs from unknown distributions.

Robustness levels Robustness can be ranked on different levels:

Level 0 No robustness measures or mitigation functionalities.

Level 1 Generalization under distribution shift. It aims at mitigating data shifts and out-of-distribution data.

Level 2 Robustness against a single risk.

Level 3 Robustness against multiple risks.

Level 4 Universal robustness against all known risks.

Level 5 Level 4 system with human-aligned and augmented robustness.

AI safety Build a system less vulnerable to adversarial attacks. This can be achieved by:

- Identifying anomalies.
- Defining safety objectives.

Reproducibility Build a system that exhibits the same behavior under the same conditions.

Remark (Repeatability vs replicability vs reproducibility).

Repeatability The same team can repeat the results under the same experimental setup.

Replicability A different team can repeat the results under the same experimental setup.

Reproducibility A different team can repeat the results with some tolerance under a different experimental setup.

Robustness requirements Two aspects have to be considered for robustness:

Robustness
requirements

Performance Capability of a model to perform a task reasonably well (humans can be used as baseline).

Vulnerability Resistance of the model to attacks. Possible sources of attack are: data poisoning, adversarial examples, flaws in the model.

Robustness approaches Robustness can be imposed with different methods at different moments of the lifecycle of the system:

Robustness
approaches

- Data sanitization,
- Robust learning,
- Extensive testing,
- Formal verification.

2.1 Robust learning

Robust learning Learn a model that is general enough to handle slightly out-of-distribution data.

Remark. It is impossible (and possibly unwanted) to have a system that models everything.

Theorem 2.1.1 (Fundamental theorem of machine learning).

$$\text{error rate} = \frac{\text{model complexity}}{\text{sample size}}$$

Corollary 2.1.1.1. If the sample size is small, the model should be simple.

Remark (Uncertainty in AI). Knowledge in AI can be divided into:

Known knowns Well-established and understood areas of research:

- Theorem proving.
- Planning in deterministic and fully-observed worlds.
- Games of perfect information.

Known unknowns Areas whose understanding is incomplete:

- Probabilistic graphical models to represent and reason on uncertainty in complex systems.
- Probabilistic machine learning that is able to quantify uncertainty.
- Planning in Markov decision problems for decision-making under uncertainty.

- Computational game theory to analyze and solve games.

Unknown unknowns Areas that we do not know are unknown. They are the natural step toward robust AI.

Remark. Robustness in biology is achieved by means of a diverse and redundant population of individuals.

2.1.1 Robustness to model errors

Robust optimization Handle uncertainty and variability through optimization methods: Robust optimization

- Assign ranges to parameters to account for uncertainty.
- Optimize in a max-min formulation aiming at maximizing the performance of the worst-case.

Remark. There is a trade-off between optimality and robustness.

Model regularization Add a penalty term to the training loss to encourage simple models. Model regularization

Theorem 2.1.2. Regularization can be interpreted as robust optimization.

Optimize risk-sensitive objectives Consider, when optimizing a reward, the variability and uncertainty associated to it (e.g., minimize variance of rewards). Optimize risk-sensitive objectives

Robust inference Deal with uncertainty, noise, or variability at inference time. Robust inference

2.1.2 Robustness to unmodeled phenomena

Model expansion Expand the models with a knowledge base. Model expansion

Remark. New knowledge might contain errors or not improve the model at all.

Causal models Define causal relations. Causal models

Portfolio of models Have multiple solvers available and use a selection method to choose the most suited in any situation. Portfolio of models

Remark. Ideally, given an instance, there should be at least a solver that performs well on it.

Anomaly detection Detect instances that deviate from the expected distribution. Anomaly detection