Natural Language Processing

Last update: 11 October 2024

Contents

1	Basi	ic text processing
	1.1	Regular expressions
		1.1.1 Basic operators
	1.2	Tokenization
		1.2.1 Data-driven tokenization
	1.3	Normalization
	1.4	Edit distance
2	Lang	guage models
	2.1	Spelling correction
	2.2	Language models
	2.3	Metrics
		2.3.1 Extrinsic evaluation
		2.3.2 Intrinsic evaluation
	2.4	N-gram model problems
		2.4.1 Overfitting
		2.4.2 Out-of-vocabulary tokens
		2.4.3 Unseen sequences
3	Text	t classification 12
	3.1	Common tasks
	3.2	Classification
	3.3	Naive Bayes

1 Basic text processing

Text normalization Operations such as:

Tokenization Split a sentence in tokens.

| Remark. Depending on the approach, a token is not always a word.

Lemmatization/stemming Convert words to their canonical form.

| Example. $\{$ sang, sung, sings $\} \mapsto$ sing

Sentence segmentation Split a text in sentences.

| Remark. A period does not always signal the end of a sentence.

Tokenization

Lemmatization/stemming

Sentence segmentation

1.1 Regular expressions

Regular expression (regex) Formal language to describe string patterns.

Regular expression (regex)

1.1.1 Basic operators

Disjunction (brackets) Match a single character between square brackets [].

| Example. / [wW] oodchuck/ matches Woodchuck and woodchuck.

Range Match a single character from a range of characters or digits.

Example.

- /[A-Z]/ matches a single upper case letter.
- /[a-z]/ matches a single lower case letter.
- /[0-9]/ matches a single digit.

Negation Match the negation of a pattern.

Example. /[^A-Z]/ matches a single character that is not an upper case letter.

Disjunction (pipe) Disjunction of regular expressions separated by |.

| Example. /groundhog | woodchuck/ matches groundhog and woodchuck.

Wildcards

Optional A character followed by ? can be matched optionally.

Example. /woodchucks?/ matches woodchuck and woodchucks.

Any . matches any character.

Kleene * A character followed by * can be matched zero or more times.

Kleene + A character followed by + must be matched at least once.

Counting A character followed by $\{n,m\}$ must be matched from n to m times.

Example.

- $\{n\}$ matches exactly n instances of the previous character.
- $\{n,m\}$ matches from n to m instances of the previous character.
- $\{n,\}$ matches at least n instances of the previous character.
- $\{,m\}$ matches at most m instances of the previous character.

Anchors

Start of line ^ matches only at the start of line.

| Example. $/^a$ / matches <u>a</u> but not ba.

End of line \$ matches only at the end of line.

| Example. /a\$/ matches \underline{a} but not \underline{a} b.

Word boundary \b matches a word boundary character.

Word non-boundary \B matches a word non-boundary character.

Aliases

- \d matches a single digit (same as [0-9]).
- \D matches a single non-digit (same as [^\d]).
- \w matches a single alphanumeric or underscore character (same as [a-zA-Z0-9_]).
- $\$ matches a single non-alphanumeric and non-underscore character (same as $[^\w]$).
- \s matches a single whitespace (space or tab).
- \S matches a single non-whitespace.

Capture group Operator to refer to previously matched substrings.

Example. In the regex /the (.*)er they were, the \left\1er they will be/, $\$ should match the same content matched by (.*).

1.2 Tokenization

Lemma Words with the same stem and roughly the same semantic meaning.

Lemma

Example. cat and cats are the same lemma.

Wordform Orthographic appearance of a word.

Wordform

Example. cat and cats do not have the same wordform.

Vocabulary Collection of text elements, each indexed by an integer.

Vocabulary

Remark. To reduce the size of a vocabulary, words can be reduced to lemmas.

Type / Wordtype Element of a vocabulary (i.e., wordforms in the vocabulary).

Type / Wordtype

Token Instance of a type in a text.

Token

Genre Topic of a text corpus (e.g., short social media comments, books, Wikipedia pages, ...).

2

Genre

Remark (Herdan's law). Given a corpus with N tokens, a vocabulary V over that corpus roughly have size:

$$|V| = kN^{\beta}$$

where the typical values are $10 \le k \le 100$ and $0.4 \le \beta \le 0.6$.

Stopwords Frequent words that can be dropped.

Stopwords

Remark. If semantics is important, stopwords should be kept. LLMs keep stopwords.

Rule-based tokenization Hand-defined rules for tokenization.

Rule-based tokenization

| Remark. For speed, simple tokenizers use regex.

Data-driven tokenization Determine frequent tokens from a large text corpus.

Data-driven tokenization

1.2.1 Data-driven tokenization

Tokenization is done by two components:

Token learner Learns a vocabulary from a given corpus (i.e., training).

Token learner

Token segmenter Segments a given input into tokens based on a vocabulary (i.e., inference).

Token segmenter

Byte-pair encoding (BPE) Based on the most frequent n-grams.

Byte-pair encoding (BPE)

Token learner Given a training corpus C, BPE determines the vocabulary as follows:

- 1. Start with a vocabulary V containing all the 1-grams of C and an empty set of merge rules M.
- 2. While the desired size of the vocabulary has not been reached:
 - a) Determine the pair of tokens $t_1 \in V$ and $t_2 \in V$ such that, among all the possible pairs, the *n*-gram $t_1 + t_2 = t_1 t_2$ obtained by merging them is the most frequent in the corpus C.
 - b) Add t_1t_2 to V and the merge rule $t_1 + t_2$ to M.

Example. Given the following corpus:

Occurrences	Tokens	
5	1 o w \$	
2	lower\$	
6	newest\$	
6	widest\$	

The initial vocabulary is: $V = \{\$, 1, o, w, e, r, n, w, s, t, i, d\}$.

At the first iteration, e + s = es is the most frequent *n*-gram. Corpus and vocabulary are updated as:

Occurrences	Tokens
5	1 o w \$
2	lower\$
6	newest\$
6	widest\$

$$V = \{\$, \texttt{l}, \texttt{o}, \texttt{w}, \texttt{e}, \texttt{r}, \texttt{n}, \texttt{w}, \texttt{s}, \texttt{t}, \texttt{i}, \texttt{d}\} \cup \{\texttt{es}\}$$

At the second iteration, es + t = est is the most frequent n-gram:

Occurrences	Tokens
5	1 o w \$
2	lower\$
6	n e w est \$
6	w i d est \$

$$V = \{\$, \texttt{l}, \texttt{o}, \texttt{w}, \texttt{e}, \texttt{r}, \texttt{n}, \texttt{w}, \texttt{s}, \texttt{t}, \texttt{i}, \texttt{d}, \texttt{es}\} \cup \{\texttt{est}\}$$

And so on...

Token segmenter Given the vocabulary V and the merge rules M, the BPE segmenter does the following:

- 1. Split the input into 1-grams.
- 2. Iteratively scan the input and do the following:
 - a) Apply a merge rule if possible.
 - b) If no merge rules can be applied, lookup the (sub)word in the vocabulary. Tokens out-of-vocabulary are marked with a special unknown token [UNK].

WordPiece Similar to BPE with the addition of merge rules ranking and a special leading/tailing set of characters (usually ##) to identify subwords (e.g., new##, ##est are possible tokens).

WordPiece

Unigram Starts with a big vocabulary and remove tokens following a loss function.

Unigram

1.3 Normalization

Normalization Convert tokens into a standard form.

Normalization

Example. U.S.A. and USA should be encoded using the same index.

Case folding Map every token to upper/lower case.

Case folding

| Remark. Depending on the task, casing might be important (e.g., US vs us).

Lemmatization Reduce inflections and variant forms to their base form.

Lemmatization

|Example. $\{$ am, are, is $\} \mapsto$ be

| Remark. Accurate lemmatization requires complete morphological parsing.

Stemming Reduce terms to their stem.

Stemming

| Remark. Stemming is a simpler approach to lemmatization.

Porter stemmer Simple stemmer based on cascading rewrite rules.

| Example. ational \mapsto ate, ing $\mapsto \varepsilon$, sses \mapsto ss.

1.4 Edit distance

Minimum edit distance Minimum number of edit operations (insertions, deletions, and substitutions) needed to transform a string into another one.

Minimum edit distance

Remark. Dynamic programming can be used to efficiently determine the minimum edit distance.

Levenshtein distance Edit distance where:

Levenshtein distance

- Insertions cost 1;
- Deletions cost 1;
- Substitutions cost 2.

Example. The Levenshtein distance between intention and execution is 8.

2 Language models

2.1 Spelling correction

Spelling correction Spelling errors can be of two types:

Spelling correction

Non-word spelling Typos that result in non-existing words. Possible candidates can be determined though a dictionary lookup.

Real-word spelling Can be:

Typographical error Typos that result in existing words.

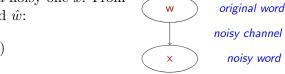
Cognitive error Due to words similarity (e.g., piece vs peace).

Noisy channel model Assumes that the observable input is a distorted form of the original word. A decoder tests word hypotheses and selects the best match.

Noisy channel model

More formally, we want a model of the channel that, similarly to Bayesian inference, determines the likelihood that a word $w \in V$ is the original word for a noisy one x. From there, we can estimate the correct word \hat{w} :



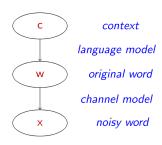


By applying (i) Bayes' rule, (ii) the fact that \hat{w} is independent of $\mathcal{P}(x)$, and (iii) that a subset $C \subseteq V$ of the vocabulary can be used, the estimate becomes:

$$\hat{w} = \arg \max_{w \in C} \underbrace{\mathcal{P}(x|w)}_{\text{channel model prior}} \underbrace{\mathcal{P}(w)}_{\text{prior}}$$

Moreover, it is reasonable to include a context c when computing the prior:

$$\hat{w} = \arg\max_{w \in C} \underbrace{\mathcal{P}(x|w)}_{\text{channel model language model}} \underbrace{\mathcal{P}(w|c)}_{\text{language model}}$$



Noisy channel spelling method Spelling correction in a noisy channel model can be done as follows:

- 1. Find candidate words with similar spelling to the input based on a distance metric (e.g., Damerau-Levenshtein which is the Levenshtein distance with the addition of adjacent transpositions).
- 2. Score each candidate based on the language and channel model:
 - Use typing features of the user.
 - Use local context.
 - Use a confusion matrix with common mistakes.

Example. Consider the sentence:

[...] was called a "stellar and versatile acress whose combination of sass and glamour has defined her [...]"

By using the Corpus of Contemporary English (COCA), we can determine the following words as candidates:

 $actress \cdot cress \cdot caress \cdot access \cdot across \cdot acres$

Language model By considering a language model without context, the priors are computed as $\mathcal{P}\left(w\right) = \frac{\mathtt{count}\left(w\right)}{|\mathtt{COCA}|}$ (where $|\mathtt{COCA}| = 404\,253\,213$):

\overline{w}	$\mathtt{count}(w)$	$\mathcal{P}\left(w\right)$
actress	9321	0.0000231
cress	220	0.000000544
caress	686	0.00000170
access	37038	0.0000916
across	120844	0.000299
acres	12874	0.0000318

Channel model By using a confusion matrix of common typos, the channel model is:

w	x w	$\mathcal{P}\left(x w\right)$
actress	cct	0.000117
cress	a #	0.00000144
caress	ac ca	0.00000164
access	r c	0.000000209
across	e o	0.0000093
acres	es e	0.0000321
acres	ssss	0.0000342

The ranking is the obtained as:

w	$\mathcal{P}\left(x w\right)\mathcal{P}\left(w\right)$
actress	$2.7\cdot 10^9$
cress	$0.00078 \cdot 10^9$
caress	$0.0028 \cdot 10^9$
access	$0.019 \cdot 10^9$
across	$2.8 \cdot 10^{9}$
acres	$1.02 \cdot 10^{9}$
acres	$1.09 \cdot 10^9$

Therefore, the most likely correction of acress for this model is across. If the previous word is considered in the context, the relevant tokens of the new language model are:

w_i	$\mathcal{P}\left(w_i w_{i-1}\right)$
actress	0.000021
across	0.000021
whose	0.001
whose	0.000006
	actress across whose

This allows to measure the likelihood of a sentence as:

 $\mathcal{P}\left(\text{versatile } \underline{\text{actress}} \text{ whose}\right) = \mathcal{P}\left(\text{actress}|\text{versatile}\right)\mathcal{P}\left(\text{whose}|\text{actress}\right) = 210 \cdot 10^{-10}$ $\mathcal{P}\left(\text{versatile } \underline{\text{across whose}}\right) = \mathcal{P}\left(\text{across}|\text{versatile}\right)\mathcal{P}\left(\text{whose}|\text{across}\right) = 1 \cdot 10^{-10}$

Finally, we have that:

 $\mathcal{P}\left(\text{versatile } \frac{\text{actress}}{\text{across}} \text{ whose} \middle| \text{versatile acress whose} \right) = 2.7 \cdot 210 \cdot 10^{-19}$ $\mathcal{P}\left(\text{versatile } \frac{\text{across}}{\text{across}} \text{ whose} \middle| \text{versatile acress whose} \right) = 2.8 \cdot 10^{-19}$

So actress is the most likely correction for acress in this model.

Remark. In practice, log-probabilities are used to avoid underflows and to make computation faster (i.e., sums instead of products).

2.2 Language models

(Probabilistic) language model Model to determine the probability of a word w in a Language model given context c:

$$\mathcal{P}\left(w|c\right)$$

Usually, it is based on counting statistics and uses as context the sequence of previous tokens:

$$\mathcal{P}\left(w_i|w_1,\ldots,w_{i-1}\right)$$

This is equivalent to computing the probability of the whole sentence, which expanded using the chain rule becomes:

$$\mathcal{P}(w_1, \dots, w_{i-1}w_i) = \mathcal{P}(w_1)\mathcal{P}(w_2|w_1)\mathcal{P}(w_3|w_{1..2})\dots\mathcal{P}(w_n|w_{1..n-1})$$
$$= \prod_{i=1}^n \mathcal{P}(w_i|w_{1..i-1})$$

Remark. Simply counting the number of occurrences of a sentence as $\mathcal{P}(w_i|w_{1..i-1}) = w_{1..i}/w_{1..i-1}$ is not ideal as there are too many possible sentences.

Markov assumption Limit the length of the context to a window of k previous tokens:

Markov assumption in language models

$$\mathcal{P}\left(w_{i}|w_{1..i-1}\right) \approx \mathcal{P}\left(w_{i}|w_{i-k..i-1}\right)$$

$$\mathcal{P}(w_{1..n}) \approx \prod_{i=1}^{n} \mathcal{P}(w_i|w_{i-k..i-1})$$

Unigram model Model without context (k = 0):

$$\mathcal{P}\left(w_{1..n}\right) \approx \prod_{i} \mathcal{P}\left(w_{i}\right)$$

Bigram model Model with a single token context (k = 1):

$$\mathcal{P}\left(w_{1..n}\right) \approx \prod_{i} \mathcal{P}\left(w_{i}|w_{i-1}\right)$$

N-gram model Model with a context of k = N - 1 tokens:

N-gram model

$$\mathcal{P}(w_{1..n}) \approx \prod_{i} \mathcal{P}(w_{i}|w_{i-N+1..i-1})$$

| Remark. N-gram models cannot capture long-range dependencies.

Estimating N-gram probabilities Consider the bigram case, the probability that a token w_i follows w_{i-1} can be determined by counting:

$$\mathcal{P}\left(w_{i}|w_{i-1}\right) = \frac{\mathtt{count}(w_{i-1}w_{i})}{\mathtt{count}(w_{i-1})}$$

| Remark. N-gram models cannot handle unknown tokens.

Remark. N-gram models capture knowledge about:

- Grammar and syntax.
- Some information about the dataset (e.g., domain, genre of corpus, cultural aspects, ...).

Generation by sampling Randomly sample tokens from the distribution of a language model.

Generation by sampling

Remark. In N-gram models $(N \ge 2)$, the distribution changes depending on the previously sampled tokens.

2.3 Metrics

2.3.1 Extrinsic evaluation

Extrinsic/downstream evaluation Compare the performance of different models on specific tasks.

Extrinsic evaluation

Remark. Extrinsic evaluation is the best approach for comparing different models, but it is often computationally expensive.

2.3.2 Intrinsic evaluation

Intrinsic evaluation Measure the quality of a model independently of the task.

Intrinsic evaluation

Perplexity (PP) Probability-based metric based on the inverse probability of a sequence (usually the test set) normalized by the number of words:

Perplexity

$$\mathcal{P}\left(w_{1..N}\right) = \prod_{i} \mathcal{P}\left(w_{i} | w_{1..i-1}\right)$$

$$PP(w_{1..N}) = \mathcal{P}\left(w_{1..N}\right)^{-\frac{1}{N}} \in [1, +\infty]$$

A lower perplexity represents a generally better model.

Example. For bigram models, perplexity is computed as:

$$\mathcal{P}\left(w_{1..N}\right) \approx \prod_{i} \mathcal{P}\left(w_{i} | w_{i-1}\right) \qquad \text{PP}(w_{1..N}) = \sqrt[N]{\prod_{i} \frac{1}{\mathcal{P}\left(w_{i} | w_{i-1}\right)}}$$

Remark (Perplexity intuition). Perplexity can be seen as a measure of surprise of a language model when evaluating a sequence.

Alternatively, it can also be seen as a weighted average branching factor (i.e., average number of possible next words that follow any word, accounting for their probabilities). For instance, consider a vocabulary of digits and a training corpus where every digit appears with uniform probability 0.1. The perplexity of any sequence using a 1-gram model is:

$$PP(w_{1..N}) = (0.1^N)^{-\frac{1}{N}} = 10$$

Now consider a training corpus where 0 occurs 91% of the time and the other digits 1% of the time. The perplexity of the sequence 0 0 0 0 0 3 0 0 0 is:

$$PP(0\ 0\ 0\ 0\ 0\ 3\ 0\ 0\ 0) = (0.91^9 \cdot 0.01)^{-\frac{1}{10}} \approx 1.73$$

Remark. Minimizing perplexity is the same as maximizing the probability of the tokens.

Remark. Perplexity can be artificially reduced by using a smaller vocabulary. Therefore, it is only reasonable to compare perplexity of models with the same vocabulary.

Remark. Perplexity is generally a bad approximation of extrinsic metrics and only works well if the test set is representative of the training data. Therefore, it is only useful to guide experiments and the final evaluation should be done through extrinsic evaluation.

2.4 N-gram model problems

2.4.1 Overfitting

N-gram models become better at modeling the training corpus for increasing values of N. This risks overfitting and does not allow to obtain a generalized model.

Overfitting

Example. A 4-gram model is able to nearly perfectly generate sentences from Shake-speare's works.

2.4.2 Out-of-vocabulary tokens

There are two types of vocabulary systems:

Closed vocabulary system All words that can occur are known.

Closed vocabulary

Open vocabulary system Unknown words are possible. They are usually handled using a dedicated token <UNK> which allows to turn an open vocabulary system into a closed one:

Open vocabulary

- Use a vocabulary and model all other words as <UNK>.
- Model infrequent words as <UNK>.

Remark. The training set must contain <UNK> tokens to estimate its distribution as it is treated as any other token.

2.4.3 Unseen sequences

Only for n-grams that occur enough times a representative probability can be estimated. For increasing values of n, the sparsity grows causing many unseen n-grams that produce a probability of 0, with the risk of performing divisions by zero (e.g., perplexity) or zeroing probabilities (e.g., when applying the chain rule).

Laplace smoothing Adds 1 to all counts and renormalizes them. Given a vocabulary V Laplace smoothing and an N-gram model, smoothing is done as follows:

$$\mathcal{P}_{\text{Laplace}}(w_i|w_{i-N+1..i-1}) = \frac{\text{count}(w_{i-N+1..i-1}w_i) + 1}{\text{count}(w_{i-N+1..i-1}) + |V|}$$

Alternatively, by only changing the numerator, it can be formulated using an adjusted count as:

$$\begin{split} \mathcal{P}_{\text{Laplace}}(w_i | w_{i-N+1..i-1}) &= \frac{c^*}{\text{count}(w_{i-N+1..i-1})} \\ c^* &= \left(\text{count}(w_{i-N+1..i-1}w_i) + 1 \right) \frac{\text{count}(w_{i-N+1..i-1})}{\text{count}(w_{i-N+1..i-1}) + |V|} \end{split}$$

where $\frac{\mathsf{count}(w_{i-N+1..i-1})}{\mathsf{count}(w_{i-N+1..i-1})+|V|}$ is a normalization factor.

Example. For a 2-gram model, Laplace smoothing is computed as:

$$\mathcal{P}_{\text{Laplace}}(w_i|w_{i-1}) = \frac{\text{count}(w_{i-1}w_i) + 1}{\text{count}(w_{i-1}) + |V|}$$

Or by using the adjusted count as:

$$\mathcal{P}_{\text{Laplace}}(w_i|w_{i-1}) = \frac{c^*}{\text{count}(w_{i-1})} \qquad c^* = \left(\text{count}(w_{i-1}w_i) + 1\right) \frac{\text{count}(w_{i-1})}{\text{count}(w_{i-1}) + |V|}$$

3 Text classification

3.1 Common tasks

Sentiment analysis/Opinion mining Detection of attitudes. It can involve detecting:

- The holder of attitude (i.e., the source).
- The target of attitude (i.e., the aspect).
- The type of attitude (e.g., positive and negative).
- The text containing the attitude.

Spam detection

Language identification

Authorship attribution

Subject category classification

3.2 Classification

Classification task Given an input x and a set of possible classes $Y = \{y_1, \dots, y_M\}$, a Classification task classifier determines the class $\hat{y} \in Y$ associated to x.

Classification can be:

Rule-based Based on fixed (possibly handwritten) rules.

Example. Blacklist, whitelist, regex, ...

In-context learning Provide a decoder (i.e., generative) large language model a In-context learning prompt describing the task and the possible classes.

Example. Zero-shot learning, few-shot learning, ...

Supervised machine learning Use a training set of N labeled data $\{(d_i, c_i)\}$ to fit a classifier.

Supervised machine learning

Rule-based

An ML model can be:

Generative Informally, it learns the distribution of the data.

Discriminative Informally, it learns to exploit the features to determine the class.

3.3 Naive Bayes

Bag-of-words (BoW) Represents a document using the frequencies of its words.

Bag-of-words (BoW)

Given a vocabulary V and a document d, the bag-of-words embedding of d is a vector in $\mathbb{N}^{|V|}$ where the i-th position contains the number of occurrences of the i-th token in d.

Sentiment analysis/Opinion mining

Multinomial naive Bayes classifier Probabilistic (i.e., generative) classifier based on the assumption that features are independent given the class.

Multinomial naive Bayes classifier

Given a document $d = \{w_1, \dots, w_n\}$, a naive Bayes classifier returns the class \hat{c} with maximum posterior probability:

$$\hat{c} = \arg \max_{c \in C} \mathcal{P}(c|d)$$

$$= \arg \max_{c \in C} \underbrace{\mathcal{P}(d|c)}_{\text{likelihood prior}} \underbrace{\mathcal{P}(c)}_{\text{prior}}$$

$$= \arg \max_{c \in C} \mathcal{P}(w_1, \dots, w_n|c) \mathcal{P}(c)$$

$$= \arg \max_{c \in C} \prod_i \mathcal{P}(w_i|c) \mathcal{P}(c)$$

$$= \arg \max_{c \in C} \sum_i \log \mathcal{P}(w_i|c) \log \mathcal{P}(c)$$

Given a training set D with N_c classes and a vocabulary V, $\mathcal{P}(w_i|c)$ and $\mathcal{P}(c)$ are determined during training by maximum likelihood estimation as follows:

$$\mathcal{P}\left(c\right) = \frac{N_c}{|D|} \qquad \mathcal{P}\left(w_i|c\right) = \frac{\mathtt{count}(w_i,c)}{\sum_{v \in V} \mathtt{count}(v,c)}$$

where count(w, c) counts the occurrences of the word w in the training samples with class c.

| Remark. Laplace smoothing is used to avoid zero probabilities.

Remark. Stop words can be removed from the training set as they are usually not relevant.

Remark. The likelihood part of the equation $(\sum_{i} \log \mathcal{P}(w_i|c))$ can be seen as a set of class-specific 1-gram language models.

Example. Given the following training set for sentiment analysis with two classes:

Class	Document
-	just plain boring
-	entirely predictable and lacks energy
-	no surprises and very few laughs
+	very powerful
+	the most fun film of the summer

We want to classify the sentence "predictable with no fun". Excluding stop words (i.e., with), we need to compute:

$$\mathcal{P}\left(+|\text{predictable with no fun}\right) = \mathcal{P}\left(+\right)\mathcal{P}\left(\text{predictable}|+\right)\mathcal{P}\left(\text{no}|+\right)\mathcal{P}\left(\text{fun}|+\right)$$

 $\mathcal{P}\left(-|\text{predictable with no fun}\right) = \mathcal{P}\left(-\right)\mathcal{P}\left(\text{predictable}|-\right)\mathcal{P}\left(\text{no}|-\right)\mathcal{P}\left(\text{fun}|-\right)$

A vocabulary of 20 tokens can be used to represent the training samples. The required likelihoods and priors with Laplace smoothing are computed as:

$$\begin{split} \mathcal{P}\left(\mathbf{+}\right) &= \frac{2}{5} \qquad \mathcal{P}\left(\mathtt{predictable}|\mathbf{+}\right) = \frac{0+1}{9+20} \quad \mathcal{P}\left(\mathtt{no}|\mathbf{+}\right) = \frac{0+1}{9+20} \quad \mathcal{P}\left(\mathtt{fun}|\mathbf{+}\right) = \frac{1+1}{9+20} \\ \mathcal{P}\left(\mathbf{-}\right) &= \frac{3}{5} \qquad \mathcal{P}\left(\mathtt{predictable}|\mathbf{-}\right) = \frac{1+1}{14+20} \quad \mathcal{P}\left(\mathtt{no}|\mathbf{-}\right) = \frac{1+1}{14+20} \quad \mathcal{P}\left(\mathtt{fun}|\mathbf{-}\right) = \frac{0+1}{14+20} \end{split}$$